

When Can I Connect a Node to the Spine?


There are certain cases where it will work just fine, and others where it can complicate life for you and your client.

Here are some visual examples of when it is OK and Not OK to connect nodes directly to the spine.

When Can I Connect a Node to the Spine?

OK: A single director class switch with no leaf switches.


- Note: A “director class switch” is made with the same 36 port chunks of silicon as 36 port leaf switches. It has 36 port spine cards, where each port is connected via copper trace on the backplane to one of the line cards. Each of the line cards is the same 36 port chunk of silicon, with half of it's ports connected inward to copper traces on the backplane (one port to each of the 18 spine cards for an sx6536) and half of it's ports visible and facing outward.



When Can I Connect a Node to the Spine?

OK: A single director class switch with one layer of leaf switches.

- Viable routing algorithms include “UpDn and Minhop” among others.
- Note: “Directly on the spine” in this case, really means on a leaf switch (line card) that has uplinks to the spine (spine cards).



When Can I Connect a Node to the Spine?

OK: A single 36 port switch based spine.

